

Assignment 5—Bayes Nets

NOTE: If you're reading this with github's Markdown renderer, you're going to see a lot of \$'s all over the place because it can't handle \LaTeX . You should read the README.pdf, instead.

Context

Up until now, the reasoning that we've done has relied on us being *certain* about the state of the world—we *knew* that Willie wasn't on stage, or our basketball team will never win, or so on. Occasionally, we ran across pieces of information that weren't explicitly, but we've assumed a closed world and the law of the excluded middle, so life hasn't been *too* complicated for us.

Unfortunately, many situations (robotics, medicine, criminal justice, economics, etc.) don't afford us with perfect knowledge of the world, nor are we granted the liberty of assuming state just because we're ignorant of some particulars. We need to be able to reason about situations where we *don't* have complete information, just a few estimates for how certain pieces fit together.

In the current scenario, you have been hired by the Wildcat-Medicine oncology department to analyze current data of patients, and help provide doctors with more accurate chances of survival. The data presented by the department is included in the `cancer_data.csv`, and includes data about a patients **Age**, **Genes**, **Cancer**, **Test**, **Treatment** and **Prognosis**. Each columns is described below: * **Age** is the given age of a patient. * **Genes** is information about whether the patient is genetically predisposed to have cancer. This is represented as a T or F value * **Cancer** is the value used to represent if the patient as Cancer or not. This is represented as a T or F value * **Test** is a data-point to see if the patient is tested positive or negative for Cancer. This is represented as N or P * **Treatment** is a data-point if the patient has received treatment yet or not. This is represented as a T or F value * **Prognosis** is the time for current patient to recover. The time will be 1, 3 or 5 years, and is represented in the data as 1, 3 or 5.

When we want to reason about uncertainty we can turn to Bayesian networks as our tool of choice. Bayes nets, in a nutshell, lets us use a bit of math to reason about the probability of a particular hypothesis being true (or false) given a set of evidence. Hypotheses are states of the world that we want to reason about, and they include things like, "does this patient have HIV?," "is there a clear path across this room?," and the infamous "am I being robbed or is this an earthquake?" Evidence is world state that we know to be one way or another, such as "I know there are no chairs here," "the patient has this set of symptoms," and so on.

We reason about hypotheses (whether or not they hold) given evidence, and for

Bayesian networks in particular manifest as questions like, “what is the probability that it isn’t raining, given that the grass is wet and it’s cloudy?”

Your Task

You have two tasks: construct the Bayesian Net from the given data and implement an `ask` function.

Constructing Bayesian Net

To construct the Bayesian Net, you will calculate the probabilities in the data and assemble the network. You are given the structure of the network, shown in the image below.

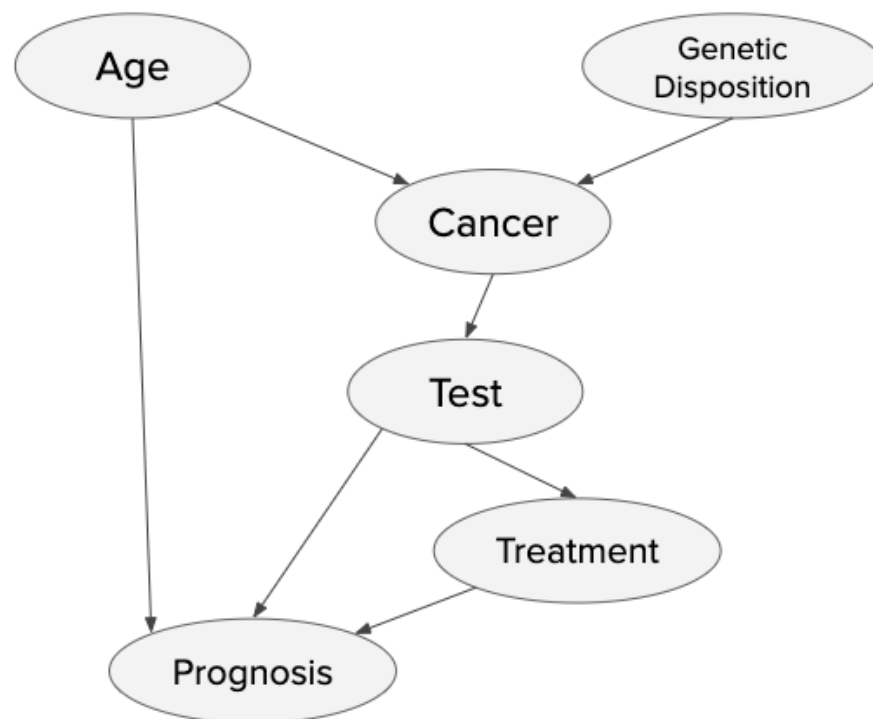


Figure 1: CancerNet

You need to calculate the conditional probability tables for each node, instantiate the node and add it to the Bayesian Net. Remember that nodes that have no incoming arrows have no conditions, and thus the probability is simply the

proportion of the data that matches the given criteria. For the other nodes, you will need to calculate the conditional probabilities from the data.

Implementing ask

Right now, `ask` doesn't do anything, but it needs to return the probability of a hypothesis given some evidence: $P(H|E)$.

The function takes four arguments: * `var` is the name of the hypothesis variable, this could be any of the data-points from the `cancer_data.csv`, i.e. `Age`, `Genes`, etc. * `value` indicates what value in the hypothesis variable are we looking for, this would be the corresponding values for each variable type, i.e. `Tests` would be N or P, `Prognosis` would be 1, 3 or 5, `Age` would be Above 55 or ≤ 55 , and `Cancer`, `Genes` and `Treatment` would be True or False. * `evidence` is the set of variables known to be True or False with a given probability. * `bn` is the `BayesNet` (in the provided module) pertaining to the problem.

To calculate $P(H|E)$, `ask` should calculate and return $\frac{P(H,E)}{\alpha}$, where:

- $P(H, E)$ is the joint probability of the hypothesis (`var = value`) and the evidence (`evidence`), and
- α is the normalization constant (the joint probability of \neg hypothesis and the evidence).

See below for more details on calculate $P(H, E)$ and α .

Testing

There provided tests in `main.py`. You should be able to easily create more test cases with this model. Additionally, you are encouraged to create your own instance of `BayesNet` that models a different domain (e.g., cavities, traffic, exam grades, or the alarm burglar example from class).

For grading the assignment we will be using the same data as in the `cancer_data.csv`.

Notes

On Joint Probability

To calculate the joint probability, we can break things down into terms that we can just lookup in the `BayesNet`. For example, the Burglary model in class does not represent $P(b, e, a)$, but it does have $P(b)$, $P(e)$, and $P(a|b, e)$.

We can handle this recursively by recognizing the following: $P(b, e, a) = P(b)P(e, a|b)$ (This is called the Chain Rule). Similarly, we know that $P(e, a|b) = P(e|b)P(a|b, e)$.

When calculating the joint probability, we need to include all the variables that may influence the final result. This includes all the variables that are parents of the variables in the call to `ask`. For example, in calculating $P(b, j, m)$, we need to include A and E (whose values we do not know; more on that in a bit). In fact, we can add in all the variables in the BayesNet, as the extra unknown variables will not affect the final result. They only add a little extra computation.

To implement this recursion you may want to introduce a new function. That function takes the list of variables in the joint probability and the collection of all known variable values. For example, for $P(e|a, b)$, the function initially takes the list of all of the variables in the bayes net as the list of variables. The known variable values can be represented with the dictionary `{'E':True, 'A': True, 'B':True}`.

This new function should handle the following conditions: * Recursion is done if there are no more variables in the list. * The next variable in the list has a known value (it is in the evidence). * The value of the next variable is not known.

In the case where the next variable has a known value, lookup the probability in the CPT using the function `probability` on the `BayesNode`. Then recurse on the rest of the variables.

If the value of the next variable is not known, we need to compute the sum of the joint probabilities when the unknown is True or False. In other words, when trying to find out $P(B, e, a)$ given unknown B and known e and a, $P(B, e, a) = P(b, e, a) + P(\neg b, e, a) = P(b)P(e, a|b) + P(\neg b)P(e, a|\neg b)$. For each of the two possibilities, lookup the probability in the CPT and recurse. Note that when recursing, we have now defined a value for the first variable, and this value needs to be included in the recursive call.

In calculating the joint probability, it is best to process the variables orderly. If the list of variables is ordered such that the parent of any variable precedes its children, then when processing the child, we will already know that value. For example, in calculating $P(a, B, e)$, if we order the variables such that we calculate $P(B, e, a)$, then when processing a, we will have specified a value for B and can thus lookup the value of $P(a|b, e)$.

To get the list of variables in order and to get the whole list of variables, use `BayesNet.variables`.

Variables and known values

In code, we represent the values of variables as a dictionary keyed on the *name* of the variable. The value of the dictionary entry is either `True` or `False` for some, `Positive` or `Negative` for `Test`, and 1, 3, or 5, for `Prognosis`.

An example:

```
{
  'Age': True,
  'Genes': False,
  'Cancer': True,
  'Test': 'Positive',
  'Treatment': False,
  'Prognosis': 3
}
```

This representation is used for the **evidence** argument in the call to **ask**. In making your recursive function call, you will want to take the given evidence and update it (in the 3rd case described above). You may notice that plenty of variables are going to be left out, and this is okay, as Bayes rule works just fine with unknowns (since that's the point...).

On The Normalization Constant

The normalization constant is the sum of: (1) the joint probability of the evidence and hypothesis and (2) the joint probability of *not* the hypothesis and the evidence. For example, the normalization constant α for $P(a|b)$ is:

$$\alpha = P(a, b) + P(\neg a, b)$$

On Making Probabilities Using the Data

In the CSV file, **cancer_data.csv**, we have exactly 10,000 patient records that are row separated. So each row is one record where the columns that are comma-separated are defined on row 0: **Age**, **Genes**, **Cancer**, **Test**, **Treatment**, **Prognosis**. This data needs to be loaded and used to calculate the probabilities used in your Bayes Net. The initial steps in processing this data has been done for you and can be found in the function **makeCancerNet()**.

For some nodes, you need to translate the raw data in **cancer_data.csv** into a more useable form. For each variable in the data that is T and F, you need to construct a BayesNode that has a sample space of **True** and **False**. For the age data, you need to split it such that all ages > 55 correspond with **True**, otherwise the age is **False**.

For example, if a record has the following:

```
Age, Genes, Cancer, Test, Treatment, Prognosis
56, T, F, N, F, 5
```

this should be interpreted as the following

```
Age, Genes, Cancer, Test, Treatment, Prognosis
True, True, False, Negative, False, 5
```

Calculating the conditional probabilities is essentially a task involving tallying. To try to simplify this task, we have provided examples on how to process the data using **pandas**. If your python installation does not already have **pandas**, you will need to install it:

```
pip install pandas
```

You can filter a pandas DataFrame by creating rules and applying the rule. The following rule can be used to get all of the records where Genes is T:

```
gen_rule = df['Genes'] == 'T'
```

You then need to apply this rule to the original DataFrame, now getting a smaller DataFrame that only has records that match the rule.

```
genes_df = df[gen_rule]
```

The length of this DataFrame will tell you how many records have Genes with a value of T.

Once you have calculated the conditional probability, you need to create the corresponding BayesNode and add it to the BayesNet.

BayesNode and BayesNet

A **BayesNet** consists of multiple **BayesNodes**. When a **BayesNode** is added to the **BayesNet**, the **BayesNet** stores the entire node in a list called **variables** and the name of the node in a list called **variable_names**. A **BayesNet** has two functions: **add** and **get_var**.

Each **BayesNode** consists of a **name** of the random variable, a list of **parents** (parent nodes in the net), sample **space** describing the set of possible values for the variable, and **cpt** (conditional probability table). A **BayesNode** has one function called **probability**, which uses its conditional probability table to calculate the probability of this **BayesNode**. For the given hypothesis and evidence, **BayesNode.probability(hypothesis, evidence)** returns the conditional probability by calculating it with the help of **cpt**.